



Popularity Prediction

Guy Kahana & Anat Peled



10/2020

Project III: Regression Models in ML

Why do we care?

◀ There are several reasons to predict a track's popularity in Spotify:

◆ **Commercial:**

- Advertisements
- Design promoting algorithms

◆ **Financial:**

- Increase revenues from advertisers
- increase commissions

◆ **Musical:**

- Engineering of a Spotify hit

The Dataset

#	Column	Non-Null	Count	Dtype
0	acousticness	169909	non-null	float64
1	artists	169909	non-null	object
2	danceability	169909	non-null	float64
3	duration_ms	169909	non-null	int64
4	energy	169909	non-null	float64
5	explicit	169909	non-null	int64
6	id	169909	non-null	object
7	instrumentalness	169909	non-null	float64
8	key	169909	non-null	int64
9	liveness	169909	non-null	float64
10	loudness	169909	non-null	float64
11	mode	169909	non-null	int64
12	name	169909	non-null	object
13	popularity	169909	non-null	int64
14	release_date	169909	non-null	object
15	speechiness	169909	non-null	float64
16	tempo	169909	non-null	float64
17	valence	169909	non-null	float64
18	year	169909	non-null	int64

dtypes: float64(9), int64(6), object(4)
memory usage: 24.6+ MB

- ◀ Spotify Dataset 1921-2020, 160k+ Tracks: [Kaggle](#), 06/2020, generated by Yamaç Eren Ay with [Spotify API](#)
- ◀ The goal: predict the popularity (0-100) of a track according to the track's features.
- ◀ > 160,000 tracks
- ◀ 19 features

Preliminary Cleaning

◀ From 19 to 16 features

```
id          : 169909 unique values
name        : 132940 unique values
artists     : 33375 unique values
release_date : 10882 unique values
year        : 100 unique values
```



◀ Remove 669 duplicates

```
1 df = df[~df.duplicated()==1]
2 df.shape

(169240, 16)
```


Overview

Correlation between numerical features (abs.)



Overview

Pairplot of Numerical Features exc. dummies



Popularity: The target


Mean: 31.66

Median: 34

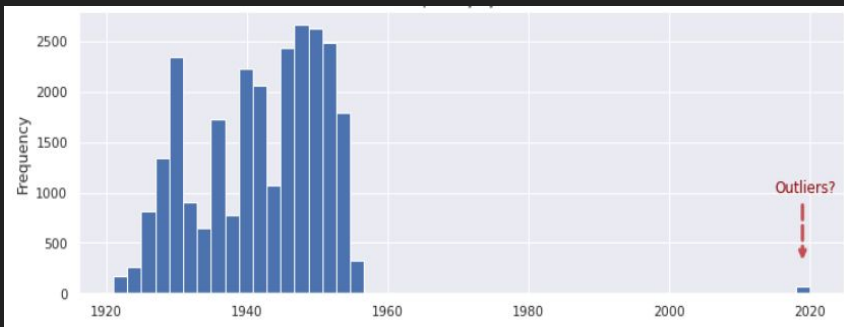
STD: 21.53

All Data

Popularity > 0



◀ Who gets 0 popularity?



```
[24] 1 df.loc[(df['popularity']==0)&(df['year']==2020), 'year'].count()
```


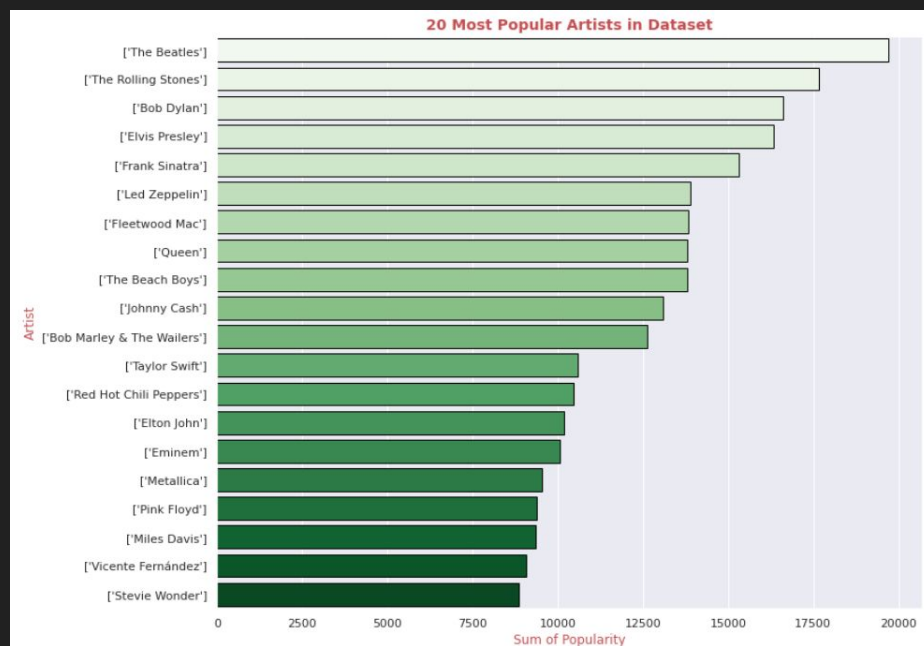

78

Categorical Features



Artists

- ◀ 33375 unique values
- ◀ Target Encoding: average the target value by category [link](#)



Artists

◀ Target Encoding:

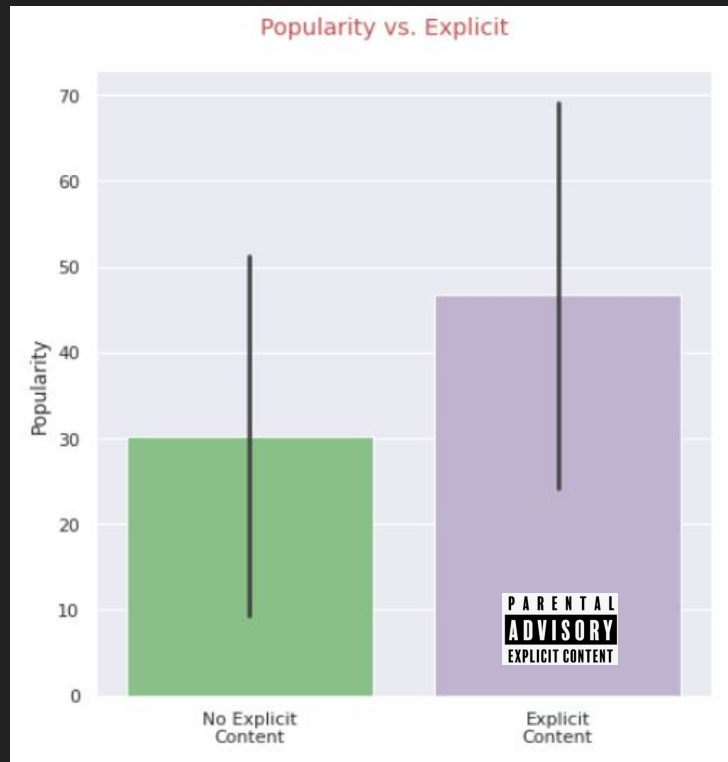
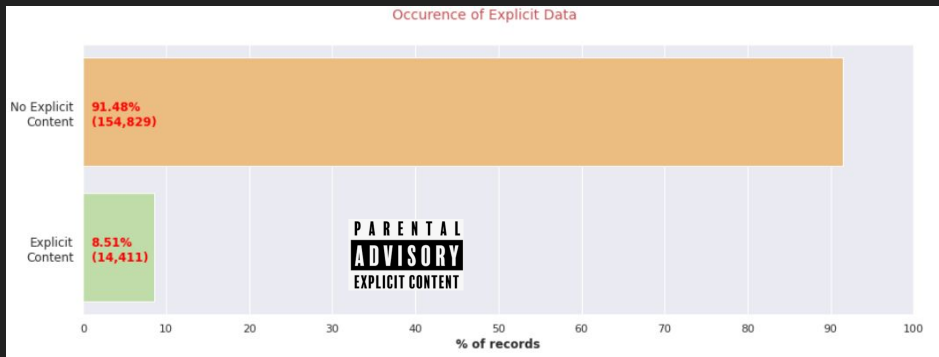
- ◆ ≤ 3 times = general mean
- ◆ > 600 times = 0
- ◆ else: unique mean

◀ MinCnt & MaxCnt set us parameters

```
1 class ArtistsTransformer():
2     """ This transformer recives a DF with a feature 'artists' of dtype object
3         and convert the feature to a float value as follows:
4         1. Replace the data with the artists mean popularity
5         2. Replace values where artists appear less than MinCnt with y.mean()
6         3. Replace values where artists appear more than MaxCnt with 0
7
8         PARAMETERS:
9         -----
10        MinCnt (int): Minimal treshold of artisits appear in dataset, default = 3
11        MaxCnt (int): Maximal treshold of artisits appear in dataset, default = 600
12
13        RERTURN:
14        -----
15        A DataFrame with converted artists str feature to ordinal floats
16    """
17
18    def __init__(self, MinCnt = 3.0, MaxCnt = 600.0):
19        self.MinCnt = MinCnt
20        self.MaxCnt = MaxCnt
21        self.artists_df = None
22
23    def fit(self, X, y):
24        self.artists_df = y.groupby(X.artists).agg(['mean', 'count'])
25        self.artists_df.loc['unknown'] = [y.mean(), 1]
26        self.artists_df.loc[self.artists_df['count'] <= self.MinCnt, 'mean'] = y.mean()
27        self.artists_df.loc[self.artists_df['count'] >= self.MaxCnt, 'mean'] = 0
28        return self
29
30    def transform(self, X, y=None):
31        X['artists'] = np.where(X['artists'].isin(self.artists_df.index), X['artists'], 'unknown')
32        X['artists'] = X['artists'].map(self.artists_df['mean'])
33        return X
```

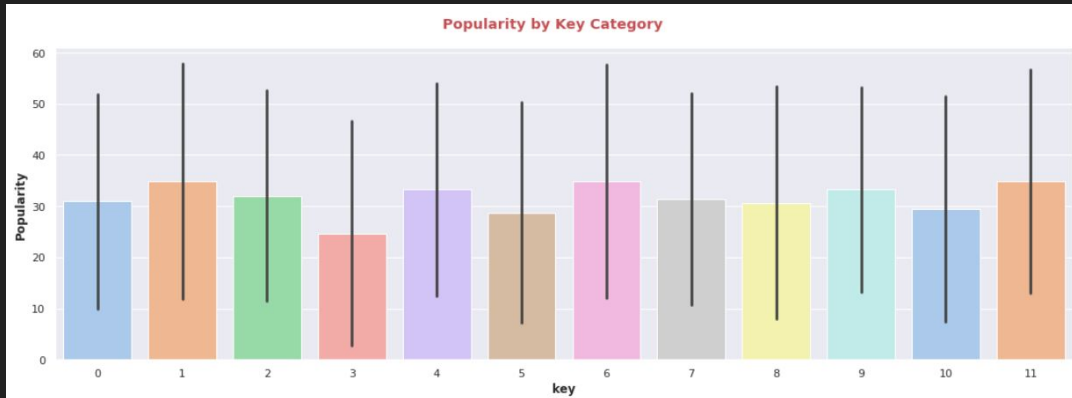
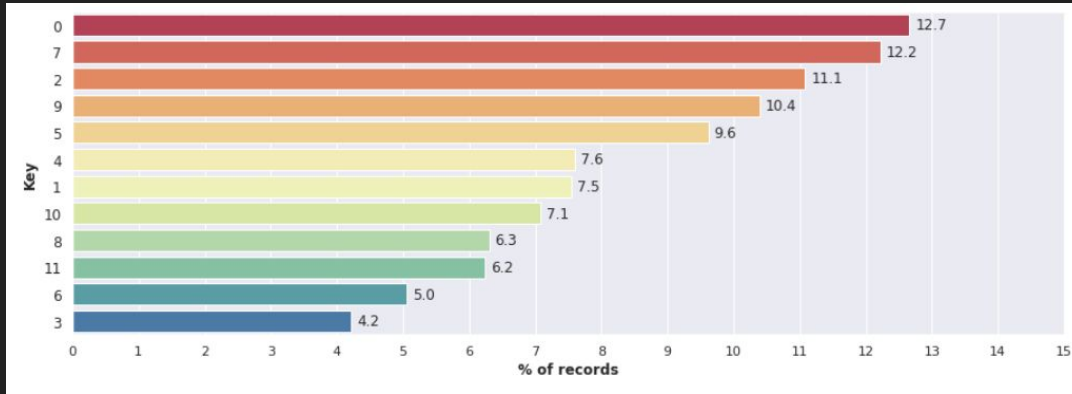
Explicit (binary)

- ◀ Warning label (abb. PAL)
- ◀ No explicit = 0, Explicit content = 1
- ◀ Linear corr. to popularity = 0.2134
- ◀ Explicit = higher mean popularity?



No Action Required

Key (0-11)

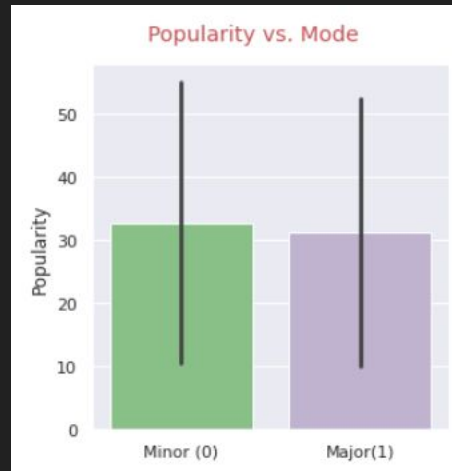



- ◀ 4-13% disposition
- ◀ Key's popularity STD is ~ 20,
- ◀ Key 3 has the lowest mean (24.67)

Mode (binary)

- ◀ Warning label (abb. PAL)
- ◀ Minor = 0, Major = 1
- ◀ Linear corr. to popularity = -0.033

No Action Required


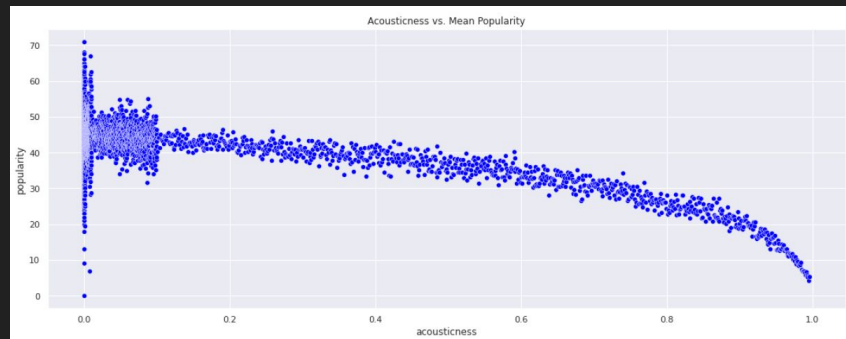


Numerical Features



Acousticness

- ◀ Higher acousticness means lower mean popularity (apply for acousticness > 0.1)
- ◀ Linear corr. to popularity = -0.5
- ◀ Majority of tracks is either close to 0 or 1

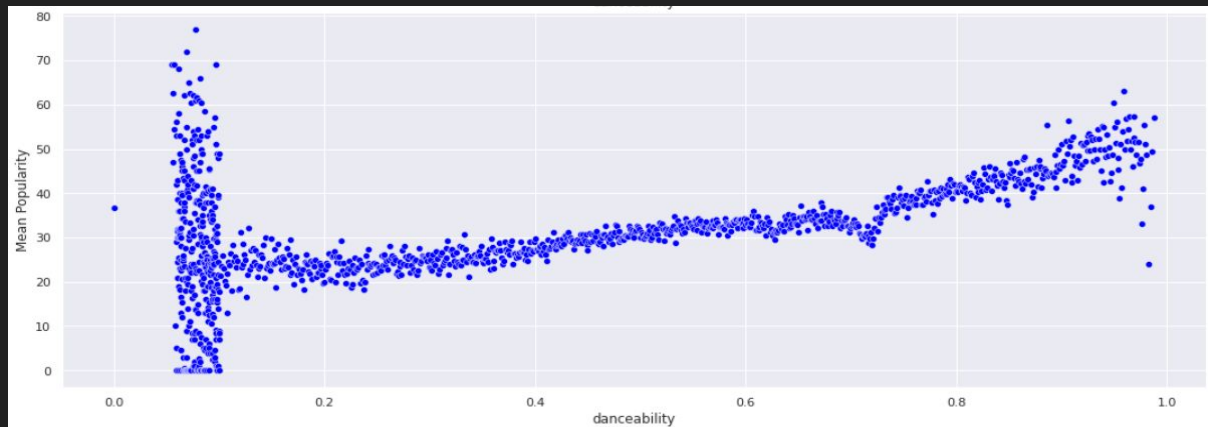
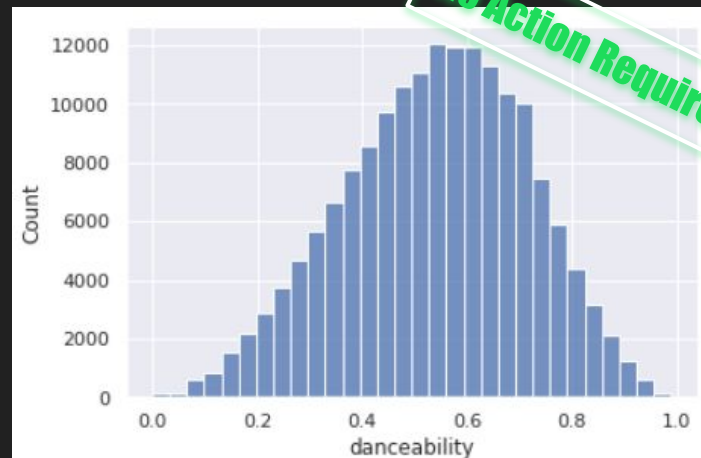


No Action Required

Danceability

- ◀ Is the song danceable based on the combination of tempo, rhythm and beat strength
- ◀ Linear corr. to popularity = 0.22

count	169240.000000
mean	0.538717
std	0.175194
min	0.000000
25%	0.418000
50%	0.549000
75%	0.668000
max	0.988000





Duration_ms

- ◀ $\text{duration_min} = \text{duration_ms} / 60,000$
- ◀ Some tracks are unexpectedly long...
- ◀ to understand the data it was divided to long tracks and short tracks:

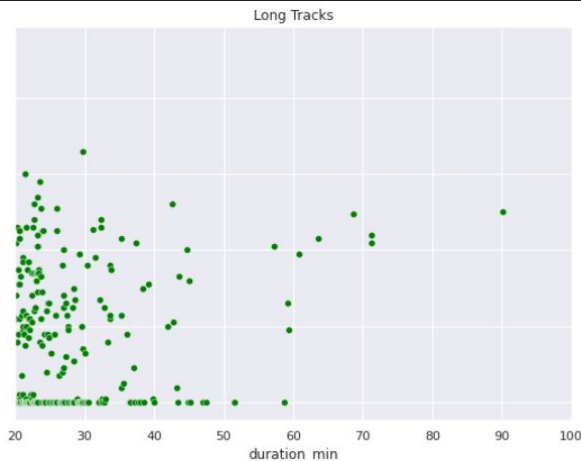

```
1 long_tracks = df.loc[df['duration_min']>20]
2 short_tracks = df.loc[df['duration_min']<=20]
```

count	169240.000000
mean	3.857319
std	2.018018
min	0.085133
25%	2.852946
50%	3.477333
75%	4.382450
max	90.058333

Duration_ms



- Long tracks are mainly classic music and atmosphere music
- There are not enough samples of tracks longer than 45 minutes

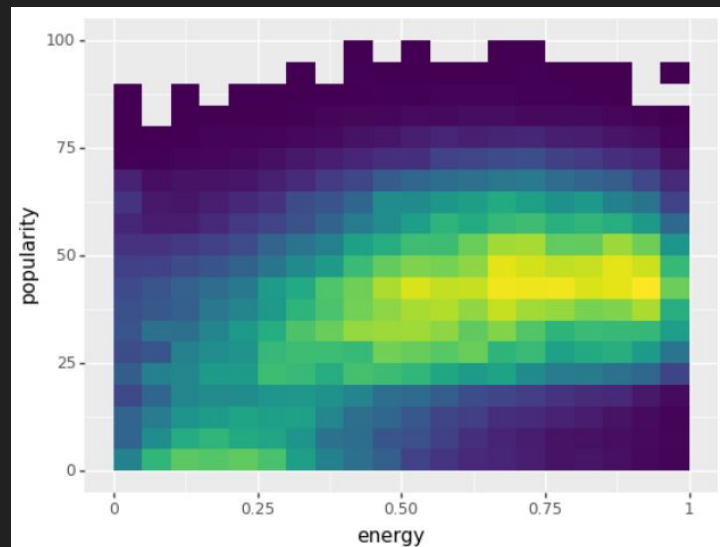
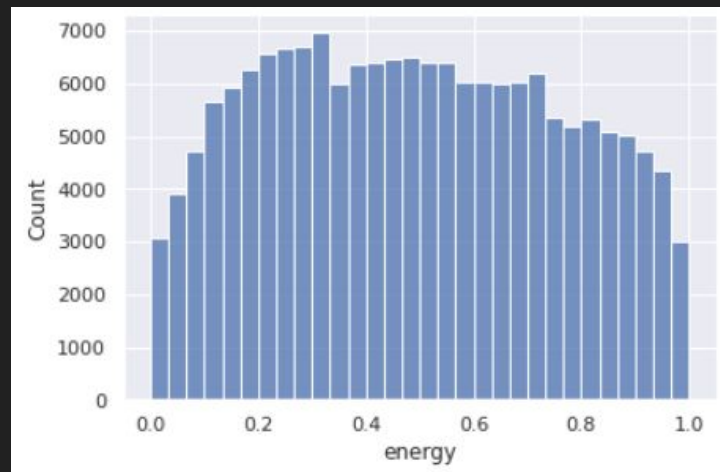


Energy

- ◀ Measure the intensity and activity, energetic track feels more fast loud and noisy
- ◀ Linear corr. to Popularity = 0.495


count	169240.000000
mean	0.489632
std	0.267099
min	0.000000
25%	0.264000
50%	0.482000
75%	0.711000
max	1.000000

No Action Required

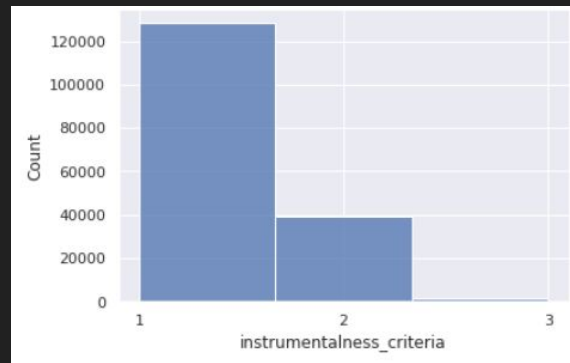



Instrumentalness

- ◀ The closer to 1 the greater likelihood the track contain NO vocals
- ◀ Linear corr. to popularity = -0.29





```
criteria= [df['instrumentalness'].between(0, 0.1),  
           df['instrumentalness'].between(0.1000001, 0.95),  
           df['instrumentalness'].between(0.950001, 1)]  
values = [1, 2, 3]  
df['instrumentalness_criteria'] = np.select(criteria, values, 0)
```



Liveness

- ◀ Detect the presence of audience, high liveness represent that the track was live
- ◀ Linear corr. to popularity = -0.076


No Action Required



Loudness


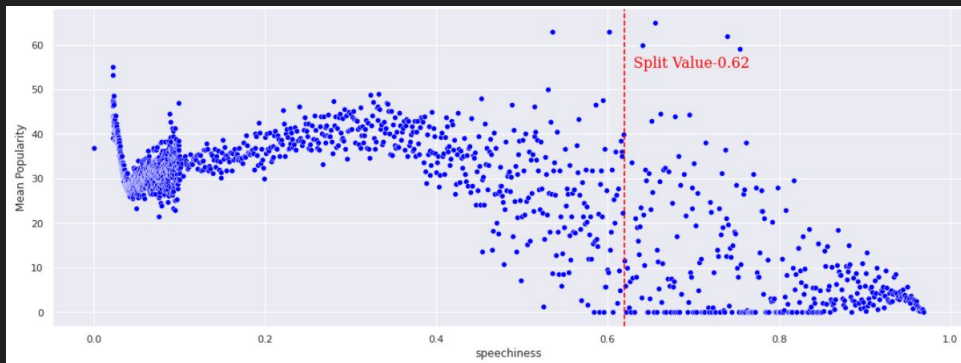

- ◀ Loudness is the quality of a sound that is the primary psychological correlate of amplitude.
- ◀ Linear corr. to popularity = 0.463

count	169240.000000
mean	-11.342179
std	5.642759
min	-60.000000
25%	-14.430000
50%	-10.453000
75%	-7.109000
max	3.855000



Speechiness



- ◀ Due to an obvious change in in speechiness trend vs. popularity, we decided to divide the records into two sub-groups.
- ◀ The cut-off point is based on a decision tree model.
- ◀ Linear corr. to popularity = -0.14



RMSE Train = 21.15
RMSE Test = 21.00

Tempo

- ◀ The speed or pace of a given piece. Derives directly from the average beat duration.
- ◀ Linear corr. to popularity = 0.13



Tempo



◀ Self made imputer with median or mean

```
1 class ReplaceZeroTransformer():
2     """Eliminates Zero values from tempo columns and replace it
3         with the median or mean of non-zero values as specified.
4         default is set to 'median'.
5     """
6
7     def __init__(self, method='median'):
8         self.method = method
9
10    def transform(self, X):
11        if self.method == 'median':
12            X.loc[X['tempo']==0, 'tempo'] = X.loc[X['tempo']>0, 'tempo'].median()
13        elif self.method == 'mean':
14            X.loc[X['tempo']==0, 'tempo'] = X.loc[X['tempo']>0, 'tempo'].mean()
15        else:
16            raise Exception("Method can be 'median' or 'mean' only!")
17        return X
```


Valence

No Action Required


- ◀ A measure describing the musical positiveness conveyed by a track
- ◀ Linear corr. = 0.01



Year



- ◀ Linear corr. to popularity = **0.88**
- ◀ Most of the years are 2000 or 4000, due to the 2000 maximal batch limit in the Spotify's API.



Pre- Processing



Pre-Processing

- ◀ Import data with relevant columns only
- ◀ Remove duplicates
- ◀ Split data to train & test
- ◀ For KNN only: sample data (frac = 0.3)

```
1 # Read column names from file
2 cols = list(pd.read_csv('data.csv', nrows =1))
3 df = pd.read_csv('data.csv', usecols=[i for i in cols if i not in ['id','name','release_date','year']])
4
5 # Remove duplicated
6 df = df[~df.duplicated()==1]
7 df = df.sample(frac=0.3)
8 #Split the data to train and test
9 X_train, X_test, y_train, y_test = split(df.drop('popularity', axis=1), df['popularity'], test_size = 0.33, random_state = 12345)
```

Transformers

- ◀ Unique feature transformers
- ◀ OneHotEncoder
- ◀ MinMaxScaling
- ◀ Target Scaling

Unique Feature Transformers

```
1 # Apply ArtistsTransformer on train and test separately
2 artists_transformer = ArtistsTransformer(MinCnt=2)
3 X_train = artists_transformer.fit(X_train, y_train).transform(X_train, y_train)
4 X_test = artists_transformer.transform(X_test, y_test)
```

```
1 def instrumentality_criteria(X):
2     X['instrumentality'] = list(map((lambda x: 1 if x < 0.1 else (3 if x > 0.95 else 2)), X.instrumentality))
3
4 instrumentality_transformer = FunctionTransformer(instrumentality_criteria)
5 instrumentality_transformer.transform(X_train)
6 instrumentality_transformer.transform(X_test)
```

Unique Feature Transformers

```
1 def speech_criteria(X):
2     X['speechiness'] = list(map((lambda x: 1 if x > 0.62 else 0), X.speechiness))
3
4 speech_transformer = FunctionTransformer(speech_criteria)
5 speech_transformer.transform(X_train)
6 speech_transformer.transform(X_test)
```

```
1 tempo_transformer = ReplaceZeroTransformer()
2 X_train = tempo_transformer.transform(X_train, 'median')
3 X_test = tempo_transformer.transform(X_test, 'median')
```


OneHotEncoder

```
1  ohe = OneHotEncoder(categories='auto', drop='first')
2
3  # Train
4  feature_arr = ohe.fit_transform(X_train[['instrumentalness','key']]).toarray()
5  columns_key = ['key_'+str(i) for i in list(set(X_train['key'].values))[1:]]
6  instrumentalness_key = ['ins_'+str(i) for i in list(set(X_train['instrumentalness'].values))[1:]]
7  feature_labels = columns_key + instrumentalness_key
8  feature_labels = np.concatenate((feature_labels), axis=None)
9  features = pd.DataFrame(feature_arr, columns = feature_labels, index = X_train.index)
10 X_train = pd.concat([X_train, features], axis=1).drop(['key','instrumentalness'], axis=1)
11
12 # Test
13 feature_arr = ohe.fit_transform(X_test[['instrumentalness','key']]).toarray()
14 columns_key = ['key_'+str(i) for i in list(set(X_test['key'].values))[1:]]
15 instrumentalness_key = ['ins_'+str(i) for i in list(set(X_test['instrumentalness'].values))[1:]]
16 feature_labels = columns_key + instrumentalness_key
17 feature_labels = np.concatenate((feature_labels), axis=None)
18 features = pd.DataFrame(feature_arr, columns = feature_labels, index = X_test.index)
19 X_test = pd.concat([X_test, features], axis=1).drop(['key','instrumentalness'], axis=1)
```

MinMaxScaler

```
1 scaler = MinMaxScaler()
2 cols = ['artists', 'duration_ms', 'loudness', 'tempo']
3 X_train[cols] = scaler.fit_transform(X_train[cols])
4 X_test[cols] = scaler.fit_transform(X_test[cols])
```

Target Scaling

```
1 # Divide the popularity by 100
2 y_train = y_train / 100
3 y_test = y_test / 100
```

Prepared Data

```
X_train shape is: (113216, 25)
y_train shape is: (113216,)
X_test shape is: (55764, 25)
y_test shape is: (55764,)
```

```
1 y_train.describe().drop(['count', '25%', '50%', '75%'])
```

```
mean    0.317545
std      0.215262
min      0.000000
max      0.990000
Name: popularity, dtype: float64
```

	acousticness	artists	danceability	duration_ms	energy	explicit	liveness	loudness	mode	speechiness	tempo	valence
mean	0.490510	0.376646	0.538910	0.041969	0.488863	0.086066	0.206873	0.774618	0.708813	0.024361	0.404491	0.532166
std	0.375787	0.202948	0.175158	0.022925	0.267015	0.280462	0.177242	0.089876	0.454311	0.154166	0.143558	0.261987
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	0.996000	1.000000	0.986000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	key_1	key_2	key_3	key_4	key_5	key_6	key_7	key_8	key_9	key_10	key_11	ins_2	ins_3
mean	0.232591	0.009318	0.075908	0.111398	0.042167	0.075537	0.095066	0.050187	0.123233	0.063286	0.103934	0.070600	0.062509
std	0.422485	0.096082	0.264852	0.314625	0.200971	0.264257	0.293307	0.218332	0.328706	0.243478	0.305176	0.256156	0.242078
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

Regression Models




Linear Regression

Score Train	Score Test	Validation	Parameters
0.129	0.133	96.99%	All data mincnt = 3, test size = 0.33
0.127	0.132	96.21%	All data mincnt = 3, test size = 0.2
0.206	0.206	100%	corr > 0.2, mincnt = 3
0.124	0.131	94.65%	All data, mincnt = 2, test size = 0.33
0.123	0.129	95.34%	All data, mincnt = 2, test size = 0.2


Linear Regression (test size = 0.33, MinCnt = 2)

$$y = 0.04 + -0.10acousticness + 0.62artists + 0.14danceability + 0.04duration_m + 0.03energy + 0.05explicit + -0.05liveness + 0.08loudness + -0.00mode + -0.11speechiness + 0.03tempo + -0.10valence + -0.02key_1 + 0.04key_2 + 0.01key_3 + -0.00key_4 + -0.01key_5 + 0.00key_6 + 0.00key_7 + 0.01key_8 + -0.00key_9 + 0.00key_{10} + -0.00key_{11} + -0.00ins_2 + -0.00ins_3$$



Linear Regression (test size = 0.2, MinCnt = 2)

$$y = 0.03 + -0.10acousticness + \boxed{0.63artists + 0.14danceability} + 0.05duration_ms + 0.03energy + 0.05explicit + -0.05liveness + 0.08loudness + -0.00mode + -0.10speechiness + 0.03tempo + -0.10valence + -0.02key_1 + 0.03key_2 + 0.01key_3 + -0.00key_4 + -0.01key_5 + 0.00key_6 + 0.00key_7 + 0.01key_8 + -0.00key_9 + 0.00key_{10} + -0.00key_{11} + -0.00ins_2 + -0.00ins_3$$



Linear Regression

```
1 import statsmodels.api as sm
2 model = sm.OLS(y_train, X_train).fit()
3 model.summary()
```

OLS Regression Results

Dep. Variable:	popularity	R-squared (uncentered):	0.895
Model:	OLS	Adj. R-squared (uncentered):	0.895
Method:	Least Squares	F-statistic:	3.847e+04
Date:	Thu, 15 Oct 2020	Prob (F-statistic):	0.00
Time:	07:58:03	Log-Likelihood:	75239.
No. Observations:	113216	AIC:	-1.504e+05
Df Residuals:	113191	BIC:	-1.502e+05
Df Model:	25		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
acousticness	-0.0996	0.002	-60.610	0.000	-0.103	-0.096
artists	0.6236	0.002	263.162	0.000	0.619	0.628
danceability	0.1417	0.003	49.197	0.000	0.136	0.147
duration_ms	0.0633	0.016	3.839	0.000	0.031	0.096
energy	0.0210	0.003	7.311	0.000	0.015	0.027
explicit	0.0479	0.001	32.408	0.000	0.045	0.051
liveness	-0.0513	0.002	-23.676	0.000	-0.056	-0.047
loudness	0.1246	0.004	29.688	0.000	0.116	0.133
mode	-0.0032	0.001	-3.816	0.000	-0.005	-0.002
speechiness	-0.1058	0.003	-40.035	0.000	-0.111	-0.101
tempo	0.0317	0.003	11.897	0.000	0.027	0.037
valence	-0.0999	0.002	-50.679	0.000	-0.104	-0.096
key_1	-0.0198	0.001	-20.962	0.000	-0.022	-0.018
key_2	0.0423	0.004	10.869	0.000	0.035	0.050
key_3	0.0076	0.002	4.487	0.000	0.004	0.011
key_4	-0.0011	0.002	-0.746	0.456	-0.004	0.002
key_5	-0.0062	0.002	-2.954	0.003	-0.010	-0.002
key_6	0.0036	0.002	2.106	0.035	0.000	0.007
key_7	0.0019	0.002	1.216	0.224	-0.001	0.005
key_8	0.0077	0.002	3.932	0.000	0.004	0.012
key_9	0.0014	0.001	0.984	0.325	-0.001	0.004
key_10	0.0042	0.002	2.358	0.018	0.001	0.008
key_11	-0.0007	0.002	-0.444	0.657	-0.004	0.002
ins_2	-0.0004	0.002	-0.216	0.829	-0.004	0.003
ins_3	0.0003	0.002	0.165	0.869	-0.003	0.004
Omnibus:	4560.460	Durbin-Watson:	1.993			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6698.732			
Skew:	0.390	Prob(JB):	0.00			
Kurtosis:	3.901	Cond. No.	73.6			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


KNN

Due to long running time, this model used a sampled data (frac = 0.3)

- ◀ MinCnt = 2
- ◀ 2 loops with different n_neighbors
 - ◆ 0-101, 100-201 with 5's step
- ◀ Results stored in lists
- ◀ The smallest train-test gap is n = 170

```
1  RMSE_train, RMSE_test = [], []
2
3  for i in range(100,201,5):
4      knn = KNeighborsRegressor(n_neighbors=i)
5      knn.fit(X_train,y_train)
6      y_train_pred = knn.predict(X_train)
7      knn_train_rmse = np.sqrt(mse(y_train, y_train_pred))
8      RMSE_train.append(knn_train_rmse.round(3))
9      y_test_pred = knn.predict(X_test)
10     knn_test_rmse = np.sqrt(mse(y_test, y_test_pred))
11     RMSE_test.append(knn_test_rmse.round(3))
```

KNN




Decision Tree (test size = 0.2)

Score Train	Score Test	Validation	Parameters
0.121	0.127	95.27%	max_leaf_nodes=10, mincnt = 3, test_size = 0.2
0.116	0.121	95.86%	max_leaf_nodes=15, mincnt = 3
0.114	0.119	95.80%	max_leaf_nodes=20, mincnt = 3
0.106	0.115	92.17%	max_leaf_nodes=139, mincnt = 3
0.106	0.114	92.98%	max_leaf_nodes=183, mincnt = 3
0.107	0.114	93.86%	mincnt = 2, max_leaf_nodes = 41, min_samples_split=2000
0.102	0.112	91.07%	max_leaf_nodes = 178, mincnt = 2
0.136	0.142	95.77%	max_leaf_nodes = 4, mincnt = 2

Decision Tree

(test size = 0.2)


```
1 RMSE3_train, RMSE3_test = [], []
2
3 for i in range(2,200):
4     tree = DecisionTreeRegressor(random_state = 15, max_leaf_nodes=i)
5     tree.fit(X_train, y_train)
6     y_train_pred = tree.predict(X_train).clip(0, 1)
7     train_rmse = np.sqrt(mse(y_train, y_train_pred))
8     RMSE3_train.append(train_rmse.round(3))
9     y_test_pred = tree.predict(X_test).clip(0, 1)
10    test_rmse = np.sqrt(mse(y_test, y_test_pred))
11    RMSE3_test.append(test_rmse.round(3))
```



Decision Tree

(test size = 0.2)

artists ≤ 0.288
mse = 0.046
samples = 135184
value = 0.318



Feature importances:

acousticness	: 0.087
artists	: 0.856
danceability	: 0.003
duration_ms	: 0.006
energy	: 0.003
explicit	: 0.006
liveness	: 0.004
loudness	: 0.025
mode	: 0.000
speechiness	: 0.000
tempo	: 0.001
valence	: 0.007
key_1	: 0.001
key_2	: 0.000
key_3	: 0.000
key_4	: 0.000
key_5	: 0.000
key_6	: 0.000
key_7	: 0.000
key_8	: 0.000
key_9	: 0.000
key_10	: 0.000
key_11	: 0.000
ins_2	: 0.000
ins_3	: 0.000


Decision Tree (test size = 0.33)

Score Train	Score Test	Validation	Parameters
0.124	0.129	96.12%	max_leaf_nodes=10, mincnt = 3, test_size = 0.33
0.119	0.125	95.20%	max_leaf_nodes=15, mincnt = 3
0.117	0.122	95.90%	max_leaf_nodes=20, mincnt = 3
0.108	0.116	93.10%	max_leaf_nodes=139, mincnt = 3
0.107	0.117	91.45%	max_leaf_nodes=183, mincnt = 3
0.108	0.116	93.10%	mincnt = 2, max_leaf_nodes = 41, min_samples_split=2000
0.103	0.114	90.35%	max_leaf_nodes = 180, mincnt = 2
0.138	0.143	96.50%	max_leaf_nodes = 4, mincnt = 2

Decision Tree


(test size = 0.33)

```
1 RMSE3_train, RMSE3_test = [], []
2
3 for i in range(2,200):
4     tree = DecisionTreeRegressor(random_state = 15, max_leaf_nodes=i)
5     tree.fit(X_train, y_train)
6     y_train_pred = tree.predict(X_train).clip(0, 1)
7     train_rmse = np.sqrt(mse(y_train, y_train_pred))
8     RMSE3_train.append(train_rmse.round(3))
9     y_test_pred = tree.predict(X_test).clip(0, 1)
10    test_rmse = np.sqrt(mse(y_test, y_test_pred))
11    RMSE3_test.append(test_rmse.round(3))
```



Decision Tree

(test size = 0.33)



artists ≤ 0.29
mse = 0.046
samples = 113216
value = 0.318

Feature importances:

acousticness:	0.096
artists	: 0.841
danceability:	0.004
duration_ms	: 0.007
energy	: 0.004
explicit	: 0.007
liveness	: 0.004
loudness	: 0.026
mode	: 0.000
speechiness	: 0.000
tempo	: 0.001
valence	: 0.008
key_1	: 0.001
key_2	: 0.001
key_3	: 0.000
key_4	: 0.000
key_5	: 0.000
key_6	: 0.000
key_7	: 0.000
key_8	: 0.000
key_9	: 0.000
key_10	: 0.000
key_11	: 0.000
ins_2	: 0.000
ins_3	: 0.000

Thanks. 